

Team Alpha

Compute Canada Tutorial

Accounts (*Please use your assigned accounts*):

kevuoft	Whatthefuck.12
jayluoft	Projectalpha8
leouoft	Leoleouoft@
Pololi	Pololi_5566

The following is not for the first-time setup. All the setups are finished and ready for use.

The following examples are for CPU Runtimes. However, our project requires scaled GPU runtimes (running neural nets parallelly across multiple GPUS). For GPU runtimes, pre-process the data along with EDA on jupyter notebook on CPU runtimes. Once the data and neural nets architecture is ready to run, we will take care of the job submission since it's quite complex. Don't worry about the checkpoints and pickle. We will take care of the training.

Windows:

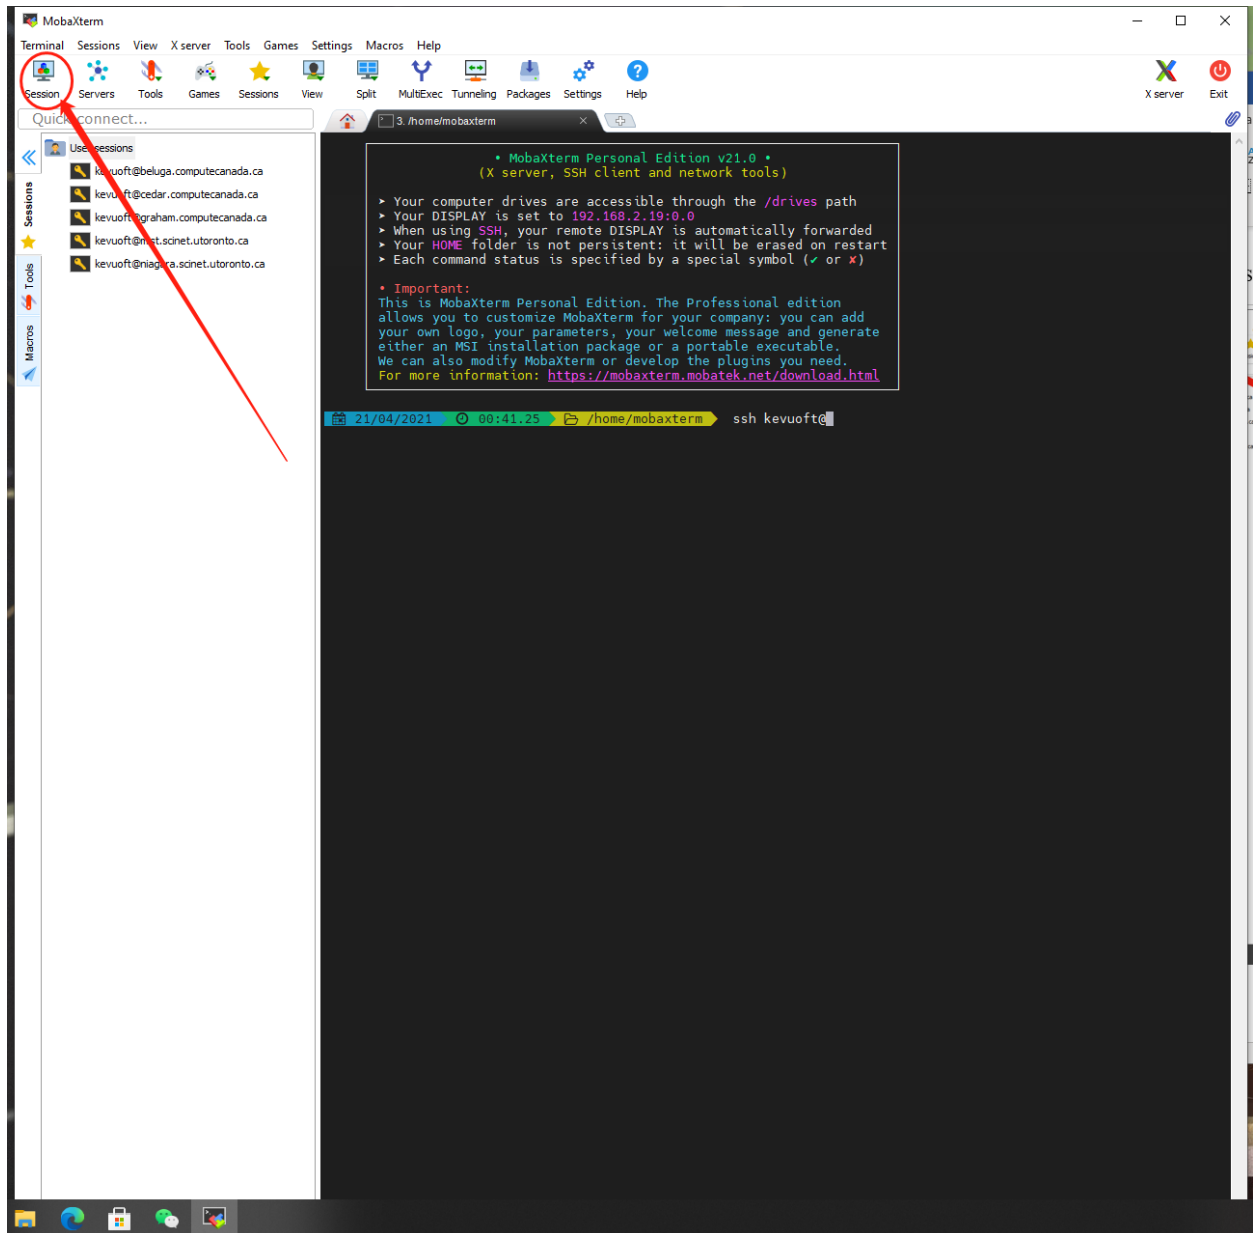
Step 1

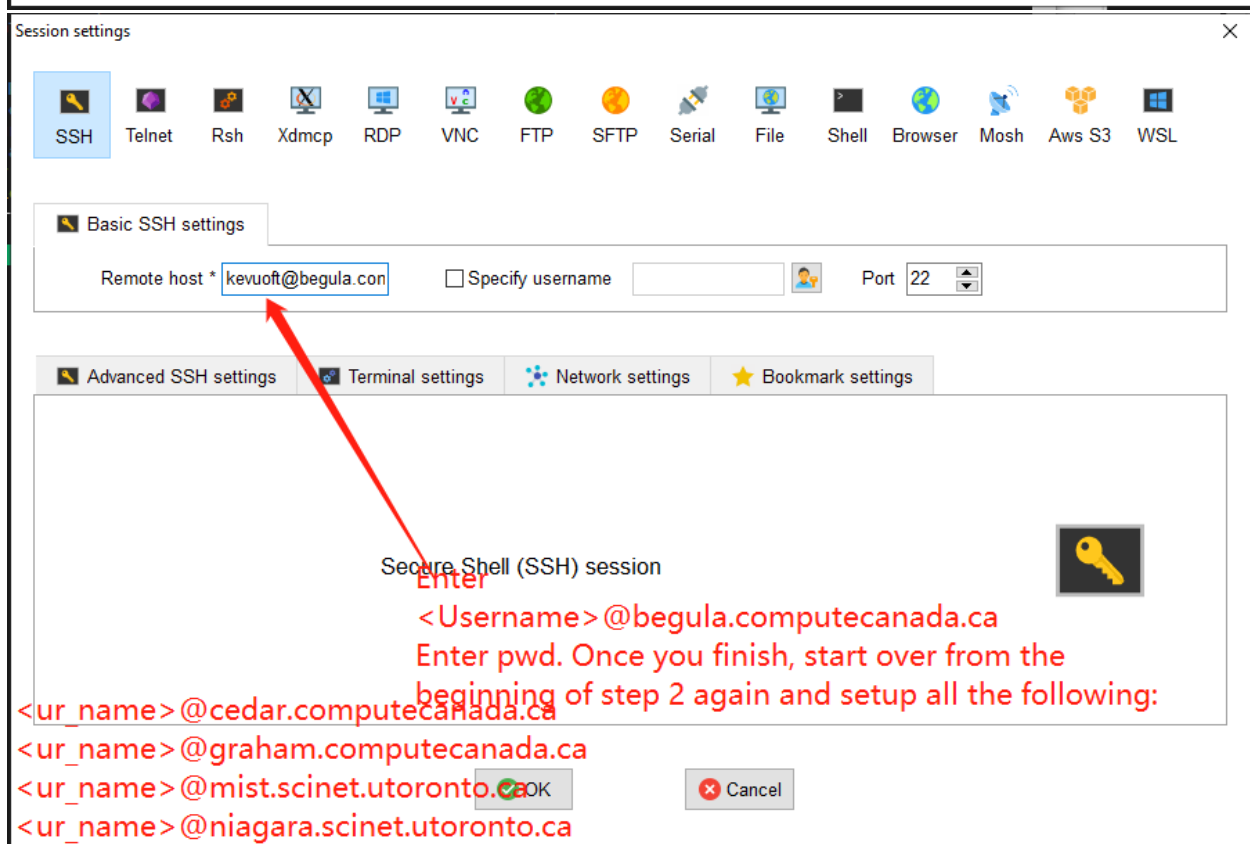
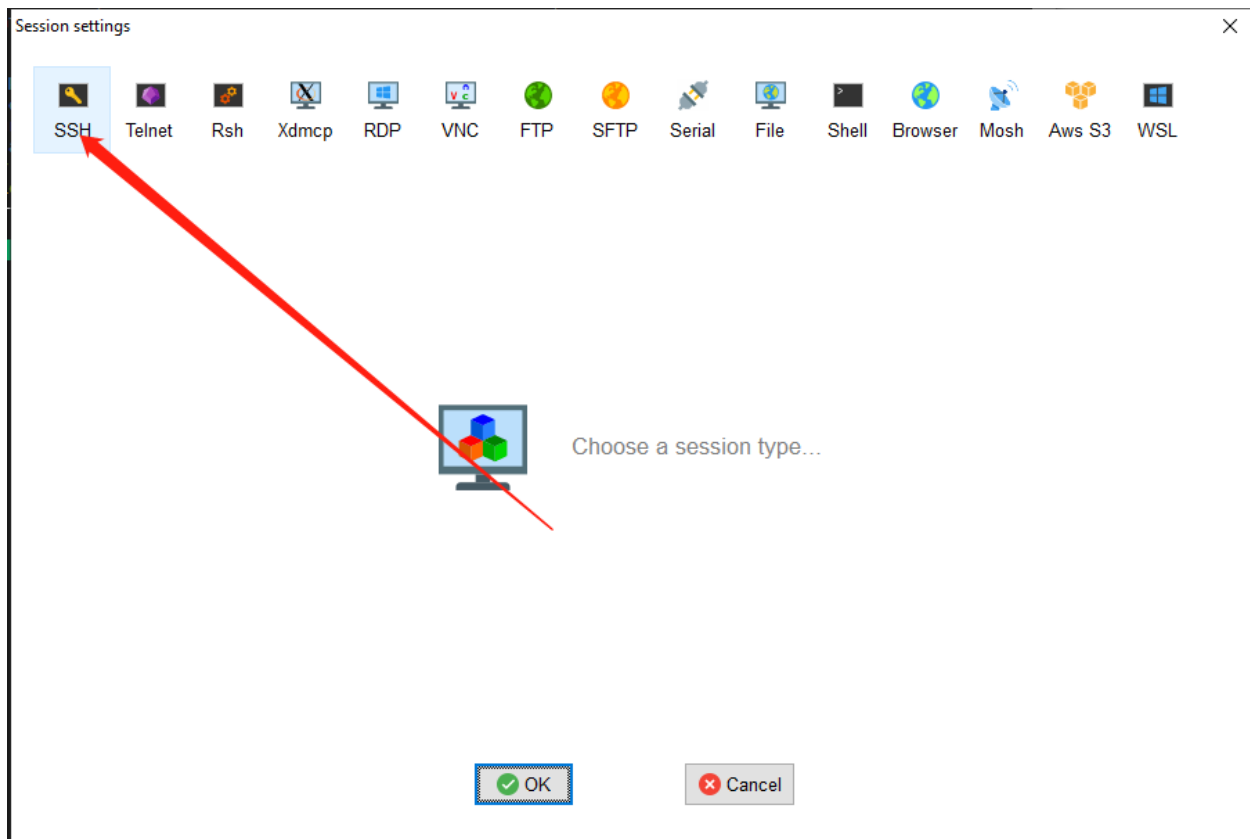
Download The MobaXterm Portable Edition

<https://mobaxterm.mobatek.net/download-home-edition.html>

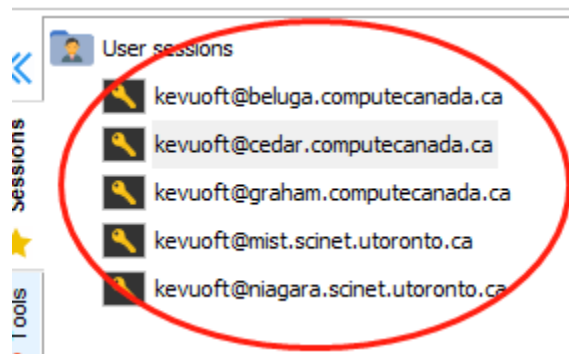
Step 2

Use the Secure Shell to connect with the clusters.





Once you have finished this step, you won't have to do it again. Instead, you would just click on the existing sessions to connect after the sessions have been set up.



The following is the connection to the compute node on the clusters. **Note that the steps are different for some of the clusters.**

Beluga & Cedar & Graham:

Step 1

Get in the directory where the data is stored.

```
-bash-4.2$ cd projects
-bash-4.2$ cd def-sekarshr
-bash-4.2$ cd Project_Alpha
-bash-4.2$
```

Step 2

Activate the virtual environment.

```
-bash-4.2$ source ~/alpha_trading/bin/activate
(alpha_trading) -bash-4.2$
```

Step 3 (Interactive)

```
(alpha_trading) -bash-4.2$ salloc --time=10:0:0 --nodes=1 --mem=200G --account=def-sekarshr srun $VIRTUAL_ENV/bin/notebook.sh
salloc: Pending job allocation 413075
salloc: job 413075 queued and waiting for resources
```

Type the command above you will get a **CPU node** with 200GB of memory and 10 hrs of runtime. Your job will be queued and you have to wait until a node becomes available for your job (This will take up to half an hour for CPU node and couple of hours for GPU nodes).

Please see the **appendix** for more commands and cluster information.

Once your job has been granted, a node will be ready for job. For interactive jobs, you will be given a token and a hostname for you to connect to your local runtime.

If you get a port error that means your localhost 8888 is probably in use. To solve this problem, run the command “jupyter notebook stop 8888” and try to request the job again.

```
(alpha_trading) -bash-4.2$ salloc --time=10:0:0 --nodes=1 --mem=180G --account=def-sekarshr srun $VIRTUAL_ENV/bin/notebook.sh
salloc: Granted job allocation 414461
[I 11:52:15.117 NotebookApp] Loading lmod extension
[I 11:52:15.285 NotebookApp] Serving notebooks from local directory: /project/6062025/Project_Alpha
[I 11:52:15.285 NotebookApp] Jupyter Notebook 6.2.0 is running at:
[I 11:52:15.285 NotebookApp] http://cdr801.int.cedar.computecanada.ca:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
[I 11:52:15.286 NotebookApp] or http://127.0.0.1:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
[I 11:52:15.286 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[C 11:52:15.332 NotebookApp]

To access the notebook, open this file in a browser:
file:///home/kevuoft/.local/share/jupyter/runtime/nbserver-31613-open.html
Or copy and paste one of these URLs:
http://cdr801.int.cedar.computecanada.ca:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
or http://127.0.0.1:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
```

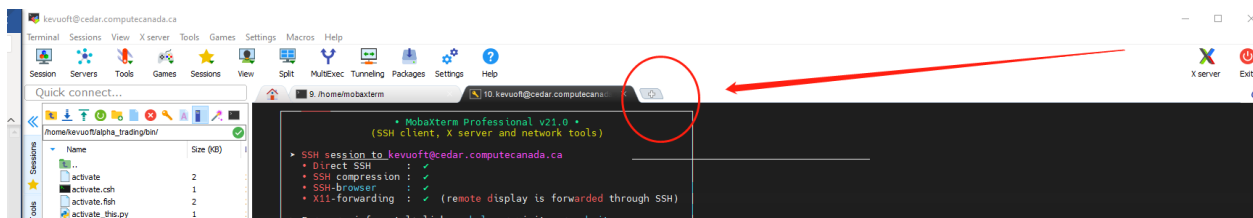
Step 4 (Interactive)

Copy the hostname and port (**Do not use control + C**, instead, select and drag on the text and right click. You will use the copy option)

```
(alpha_trading) -bash-4.2$ salloc --time=10:0:0 --nodes=1 --mem=180G --account=def-sekarshr srun $VIRTUAL_ENV/bin/notebook.sh
salloc: Granted job allocation 414461
[I 11:52:15.117 NotebookApp] Loading lmod extension
[I 11:52:15.285 NotebookApp] Serving notebooks from local directory: /project/6062025/Project_Alpha
[I 11:52:15.285 NotebookApp] Jupyter Notebook 6.2.0 is running at:
[I 11:52:15.285 NotebookApp] http://cdr801.int.cedar.computecanada.ca:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
[I 11:52:15.286 NotebookApp] or http://127.0.0.1:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
[I 11:52:15.286 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[C 11:52:15.332 NotebookApp]

To access the notebook, open this file in a browser:
file:///home/kevuoft/.local/share/jupyter/runtime/nbserver-31613-open.html
Or copy and paste one of these URLs:
http://cdr801.int.cedar.computecanada.ca:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
or http://127.0.0.1:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
```

Start a new session in MobaXterm by pressing on the plus button on top of the bash window.



Once you are in the new shell, run the command

`ssh -L 8888:<hostname:port> <username>@<cluster>.computeCanada.ca`

```
21/04/2021 15:26:18 /home/mobaxterm ssh -L 8888:cdr801.int.cedar.computeCanada.ca:8888 kevuoft@cedar.computeCanada.ca
Warning: Permanently added 'cedar.computeCanada.ca' (RSA) to the list of known hosts.
Password:
Last login: Wed Apr 21 11:31:02 2021 from bras-base-toroon213qw-grc-14-174-89-155-45.dsl.bell.ca
=====
Welcome to Cedar!

For information see: https://docs.computeCanada.ca/wiki/Cedar
Email support@computeCanada.ca for assistance and/or to report problems.
=====
-bash-4.2$
```

Once you see the cluster's welcome message, you have successfully connected to the cluster!

Step 5 (Interactive)

Go back to the previous bash window and copy the token. (Again, **do not use control + C!**)

```
(alpha_trading) -bash-4.2$ salloc --time=10:0:0 --nodes=1 --mem=180G --account=def-sekarshr srun $VIRTUAL_ENV/bin/notebook.sh
salloc: Granted job allocation 414461
[I 11:52:15.117 NotebookApp] Loading LM extension
[I 11:52:15.285 NotebookApp] Serving notebooks from local directory: /project/6062025/Project_Alpha
[I 11:52:15.285 NotebookApp] Jupyter Notebook 6.2.0 is running at:
[I 11:52:15.285 NotebookApp] http://cdr801.int.cedar.computeCanada.ca:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
[I 11:52:15.286 NotebookApp] or http://127.0.0.1:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
[I 11:52:15.286 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[C 11:52:15.332 NotebookApp]

To access the notebook, open this file in a browser:
file:///home/kevuoft/.local/share/jupyter/runtime/nbserver-31613-open.html
Or copy and paste one of these URLs:
http://cdr801.int.cedar.computeCanada.ca:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
or http://127.0.0.1:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901
```

Open your browser and open your localhost port 8888

`http://localhost:8888/?token=<token>`

`http://localhost:8888/?token=06ac9be9b9a8d970ccf433f1faa376648dce5ccb1cdb6901`

Congratulations! You have connected the jupyter notebook on your localhost.

 jupyter

Quit Logout

Files Running Clusters Softwares

Select items to perform actions on them.

0 /

	Name	Last Modified	File size
<input type="checkbox"/>	DATA	a month ago	
<input type="checkbox"/>	Factor_Bank	a month ago	
<input type="checkbox"/>	Technilal_Bank	a month ago	
<input type="checkbox"/>	MGFD25_A2_Kevin (5).ipynb	12 days ago	998 kB
<input type="checkbox"/>	A2_data.csv	a month ago	249 MB

Non-Interactive (Run <file>.py directly without interaction)

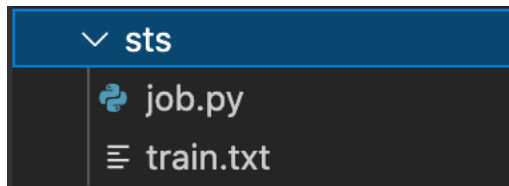
Step 1

Get in the directory where the data is stored.

```
[kevinxyc@cedar1 ~]$ cd scratch
[kevinxyc@cedar1 scratch]$ cd sts
[kevinxyc@cedar1 sts]$ █
```

This is the directory as an example in the tutorial:

The python script for executing training is “job.py” and our data is train.txt.



Step 2

Create a bash script (.sh) file in the same directory, here is a template above called myjob.sh:

```
#!/bin/bash
#SBATCH --mem=32G
#SBATCH --cpus-per-task=2
#SBATCH --time=2:0:0
#SBATCH --account=def-sekarshr
#SBATCH --gres=gpu:1
source ~/venv/bin/activate
python job.py --data train.txt
```

Line 1 stays the same.

Line 2-6 is SBATCH configuration which can be modified to your own needs (**mem** is memory required, **cpus-per-task** is number of cpus used, **time** is the total time estimated, **account** stays the same, **gres** is number of gpus used).

Line 7 is sourcing into the virtual environment you created, in this case, change **venv** to the name of your virtual env.

Line 8 is the line that execute your python script.

***If you want to output your results to an external file (say result.txt), change line 8 to be:**

```
python job.py --data train.txt > result.txt
```

Step 3

Lastly, execute the script file using sbatch and your job will be in queue for running

```
[kevinxyc@cedar1 sts]$ sbatch myjob.sh
```

After sbatching, you can use `sq` to check the status.

Once the job is done, sbatch will output some training results to `slurm-{your job id}.out` in the same directory such as `slurm-1234566.out`. We are done!

For more information, take a look at:

<https://github.com/castorini/onboarding/blob/master/docs/cc-guide.md>

Appendix I

The followings are some of the commands you may want to use for cluster Beluga, Graham and Cedar:

Command	Description	Example
salloc	This command allows you to request for interactive jobs on the cluster.	
--time=<hh:m:s>	Mandatory	--time=10:0:0 Note that your job will get shut down whether the run-job is finished or not.
--account=def-sekarshr	Mandatory	
--nodes=<#nodes>	This will specify the number of nodes that you wish to use. For interactive jobs, this will always be 1 unless using multi-node configured GPU-Accelerating jobs	--nodes=1
--mem=<mem space>	This will specify the memory that you need.	--mem=180G Note that it is the best that the memory you request does not exceeds the memory of the node that you have requested for. (Requesting for only 100G will satisfy any node)
--constraint=<device name>	This will specify the type of cpu/gpu that you would like to use.	Using only Intel Skylake cpus: --constraint=skylake Using only Intel Broadwell cpus: --constraint=broadwell Specifying the constraint may take you a long time to wait.
--gres=gpu:<device name>:<#GPUs>	Specify the gpus that you request.	--gres=gpu:v100:4 --gres=gpu:p100:4 --gres=gpu:t4:4 Only use this command on Nodes that have GPU

		For GPU jobs, please prioritize the Mist Cluster.
srun \$VIRTUAL_ENV/bin/notebook.sh	This will run the notebook.sh executable script in the bin. For interactive jobs on Begula, Cedar and Graham, you will have to always include this command at the end of the salloc request.	

Appendix II

nodes ↕	cores ↕	available memory ↕	CPU ↕	storage ↕	GPU ↕
160	40	92G or 95000M	2 x Intel Gold 6148 Skylake @ 2.4 GHz	1 x SSD 480G	-
579	40	186G or 191000M	2 x Intel Gold 6148 Skylake @ 2.4 GHz	1 x SSD 480G	-
10				6 x SSD 480G	
51	40	752G or 771000M	2 x Intel Gold 6148 Skylake @ 2.4 GHz	1 x SSD 480G	-
2				6 x SSD 480G	
172	40	186G or 191000M	2 x Intel Gold 6148 Skylake @ 2.4 GHz	1 x NVMe SSD 1.6T	4 x NVidia V100SXM2 (16G memory), connected via NVLink
2	64	4000G or 4096000M	2 x AMD EPYC 7502 Rome @ 2.5 GHz	1 x NVMe SSD 960G	-
1	32	375G or 384000M	2 x Intel Gold 6226R Cascade Lake @ 2.9 GHz	2 x SSD 480G	8 x NVidia T4 (16G memory)

- To get a larger \$SLURM_TMPDIR space, a job can be submitted with --tmp=xG, where x is a value between 350 and 2490.
- The 4 TB AMD nodes can be requested with --partition=c-slarge. Note: these nodes do not support AVX512 instructions.
- The T4 GPUs are not yet available via Slurm; only CPU cores are usable.

Cluster 1 Begula

Cedar has a total of 94,528 CPU cores for computation, and 1352 GPU devices; note that Turbo Boost is deactivated for the ensemble of Cedar nodes.

nodes ↕	cores ↕	available memory ↕	CPU ↕	storage ↕	GPU ↕
576	32	125G or 128000M	2 x Intel E5-2683 v4 Broadwell @ 2.1Ghz	2 x 480G SSD	-
96	32	250G or 257000M	2 x Intel E5-2683 v4 Broadwell @ 2.1Ghz	2 x 480G SSD	-
24	32	502G or 515000M	2 x Intel E5-2683 v4 Broadwell @ 2.1Ghz	2 x 480G SSD	-
24	32	1510G or 1547000M	2 x Intel E5-2683 v4 Broadwell @ 2.1Ghz	2 x 480G SSD	-
4	32	3022G or 3095000M	4 x Intel E7-4809 v4 Broadwell @ 2.1Ghz	2 x 480G SSD	-
114	24	125G or 128000M	2 x Intel E5-2650 v4 Broadwell @ 2.2GHz	1 x 800G SSD	4 x NVIDIA P100 Pascal (12G HBM2 memory)
32	24	250G or 257000M	2 x Intel E5-2650 v4 Broadwell @ 2.2GHz	1 x 800G SSD	4 x NVIDIA P100 Pascal (16G HBM2 memory)
192	32	187G or 192000M	2 x Intel Silver 4216 Cascade Lake @ 2.1GHz	1 x 480G SSD	4 x NVIDIA V100 Volta (32G HBM2 memory)
640	48	187G or 192000M	2 x Intel Platinum 8160F Skylake @ 2.1Ghz	2 x 480G SSD	-
768	48	187G or 192000M	2 x Intel Platinum 8260 Cascade Lake @ 2.4Ghz	2 x 480G SSD	-

Cluster 2 Cedar

Node characteristics [\[edit\]](#)

A total of 41,548 cores and 520 GPU devices, spread across 1,185 nodes of different types; note that Turbo Boost is activated for the ensemble of Graham nodes.

nodes ↕	cores ↕	available memory ↕	CPU ↕	storage ↕	GPU ↕
903	32	125G or 128000M	2 x Intel E5-2683 v4 Broadwell @ 2.1GHz	960GB SATA SSD	-
24	32	502G or 514500M	2 x Intel E5-2683 v4 Broadwell @ 2.1GHz	960GB SATA SSD	-
56	32	250G or 256500M	2 x Intel E5-2683 v4 Broadwell @ 2.1GHz	960GB SATA SSD	-
3	64	3022G or 3095000M	4 x Intel E7-4850 v4 Broadwell @ 2.1GHz	960GB SATA SSD	-
160	32	124G or 127518M	2 x Intel E5-2683 v4 Broadwell @ 2.1GHz	1.6TB NVMe SSD	2 x NVIDIA P100 Pascal (12GB HBM2 memory)
7	28	178G or 183105M	2 x Intel Xeon Gold 5120 Skylake @ 2.2GHz	4.0TB NVMe SSD	8 x NVIDIA V100 Volta (16GB HBM2 memory). Note that one node is only populated with 6 GPUs.
2	40	377G or 386048M	2 x Intel Xeon Gold 6248 Cascade Lake @ 2.5GHz	5.0TB NVMe SSD	8 x NVIDIA V100 Volta (32GB HBM2 memory), NVLINK
6	16	192G or 196608M	2 x Intel Xeon Silver 4110 Skylake @ 2.10GHz	11.0TB SATA SSD	4 x NVIDIA T4 Turing (16GB GDDR6 memory)
30	44	192G or 196608M	2 x Intel Xeon Gold 6238 Cascade Lake @ 2.10GHz	5.8TB NVMe SSD	4 x NVIDIA T4 Turing (16GB GDDR6 memory)
72	44	192G or 196608M	2 x Intel Xeon Gold 6238 Cascade Lake @ 2.10GHz	879GB SATA SSD	-

Cluster 3Graham

Node characteristics [\[edit\]](#)

- CPU: 2 sockets with 20 Intel Skylake cores (2.4GHz, AVX512), for a total of 40 cores per node
- Computational performance: 3.07 TFlops theoretical peak.
- Network connection: 100Gb/s EDR Dragonfly+
- Memory: 202 GB (188 GiB) of RAM, i.e., a bit over 4GiB per core.
- Local disk: none. GPUs/Accelerators: none.
- Operating system: Linux CentOS 7

Cluster 4Niagara

Installed	Dec 2019
Operating System	Red Hat Enterprise Linux 7.6
Number of Nodes	54 IBM AC922
Interconnect	Mellanox EDR
Ram/Node	256 GB
GPUs/Node	4 V100-SMX2-32GB
Login/Devel Node	mist.scinet.utoronto.ca
Vendor Compilers	NVCC, IBM XL
Queue Submission	Slurm

Cluster 5Mist

Appendix III

Cluster Begula, Graham and Cedar first time setup (Interactive)

```
module purge
```

```
module load openmpi
```

```
module load cuda
```

```
module load cudnn
```

```
module load nccl
```

```
module load python/3.8
```

```
virtualenv $HOME/alpha_trading
```

```
source ~/alpha_trading/bin/activate
```

```
pip install --no-index jupyter
```

```
echo -e '#!/bin/bash\nunset XDG_RUN TIME_DIR\njupyter notebook --ip $(hostname -f) --no-browser --port=8888 --NotebookApp.port_retries=0' > $VIRTUAL_ENV/bin/notebook.sh
```

```
chmod u+x $VIRTUAL_ENV/bin/notebook.sh
```

```
pip install jupyterlmod
```

```
jupyter nbextension enable --py jupyterlmod --sys-prefix
```

```
jupyter serverextension enable --py jupyterlmod --sys-prefix
```

```
pip install nbserverproxy
```

```
jupyter serverextension enable --py nbserverproxy --sys-prefix
```

```
pip install jupyter_http_over_ws
```

```
jupyter serverextension enable --py jupyter_http_over_ws
```

```
pip install --no-index tensorflow
```

```
pip install --no-index sklearn
```

```
pip install --no-index torch
```

```
pip install --no-index keras
```

```
pip install --no-index pandas
```

```
pip install --no-index numpy
```

```
pip install --no-index matplotlib
```

```
pip install --no-index plotly
```

```
pip install --no-index seaborn
```

```
pip install --no-index xgboost
```

```
pip install --no-index featuretools
```